An introduction to natural language processing in Python

Jeremy R. Manning

Computational Foundations for Neuroscience
Dartmouth College 🌲

November 2, 2023

What's covered in this tutorial?

- Today's NLP tutorial provides some background and shows how to apply several different text embedding models to a conversations dataset.
- The tutorial also shows how to implement a simple chatbot using `langchain`.
- Suggestion: open the tutorial in Google Colaboratory and select "Run All" from the "Runtime" menu, while I'm going through the slides. Then you can play around with the different pieces once everything has loaded.

What is natural language processing?

- Branch of computational linguistics
- Use computational approaches to process, analyze, and understand language

Early tasks: part of speech tagging

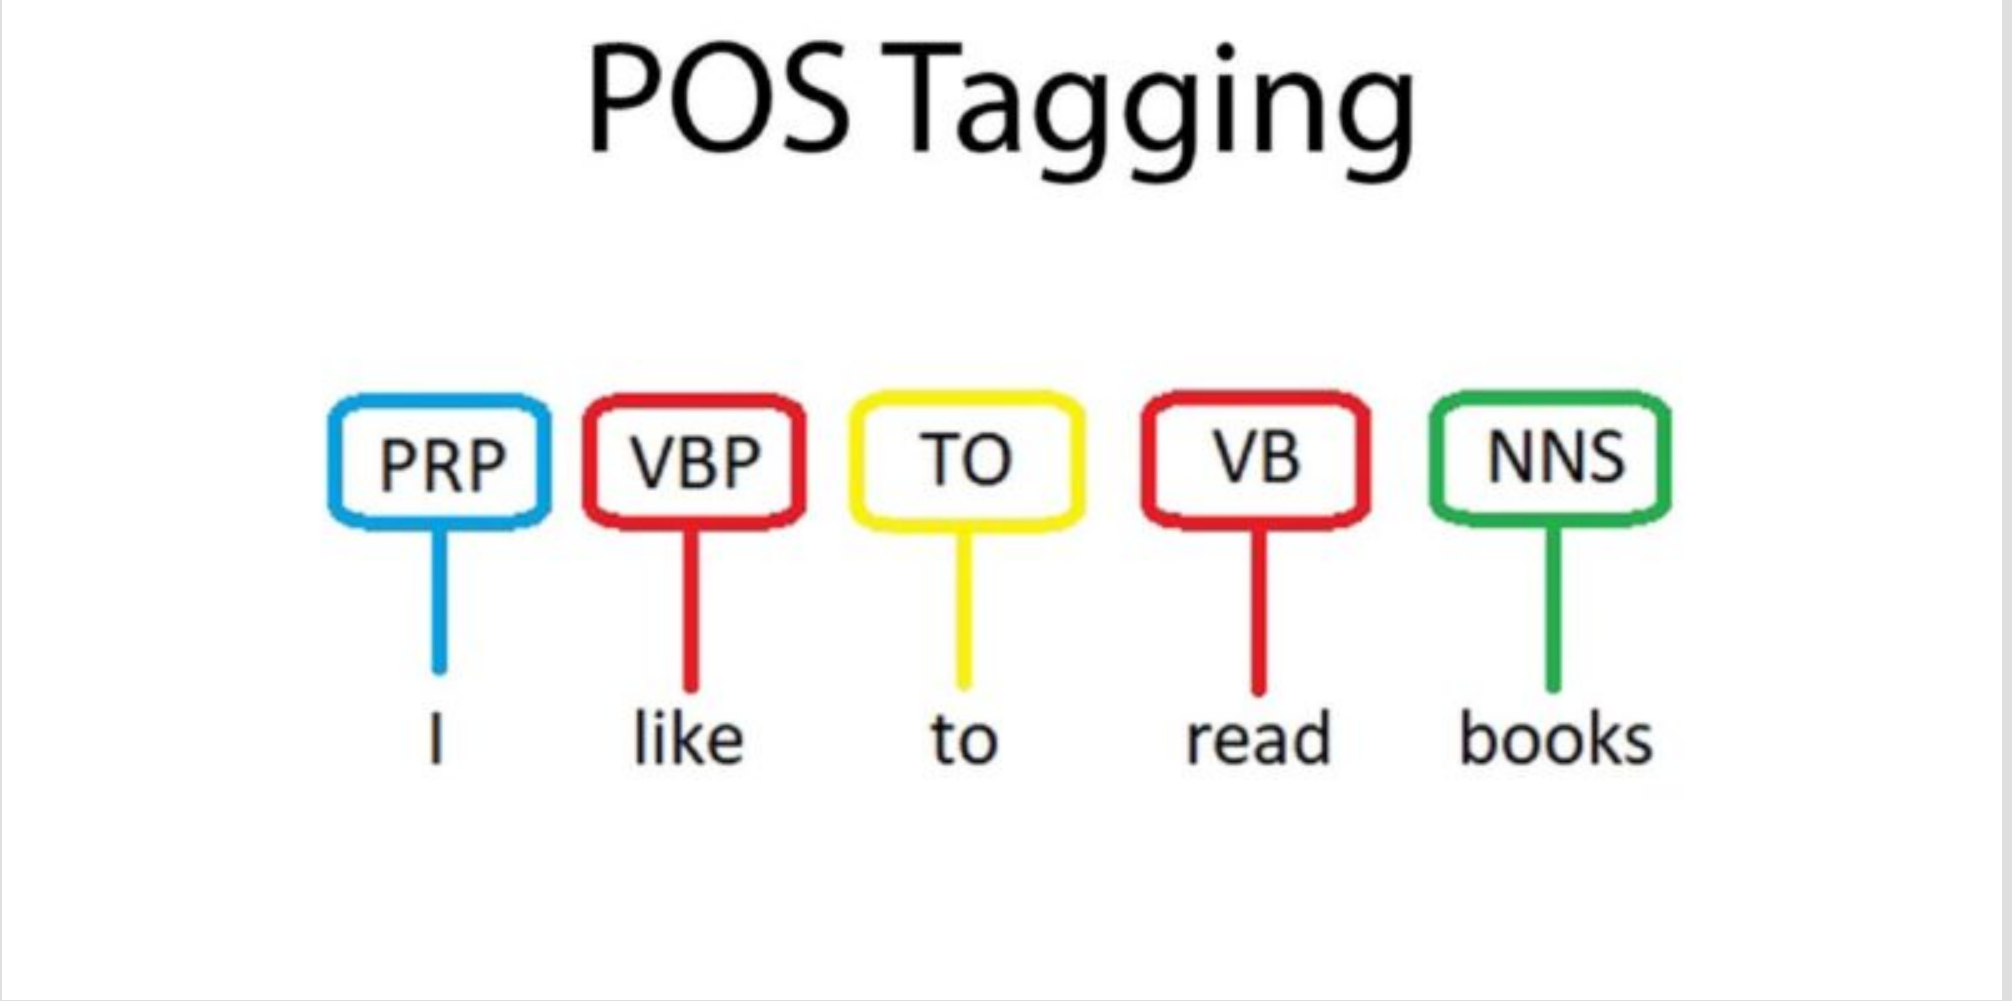

Early tasks: stemming/lemmatization

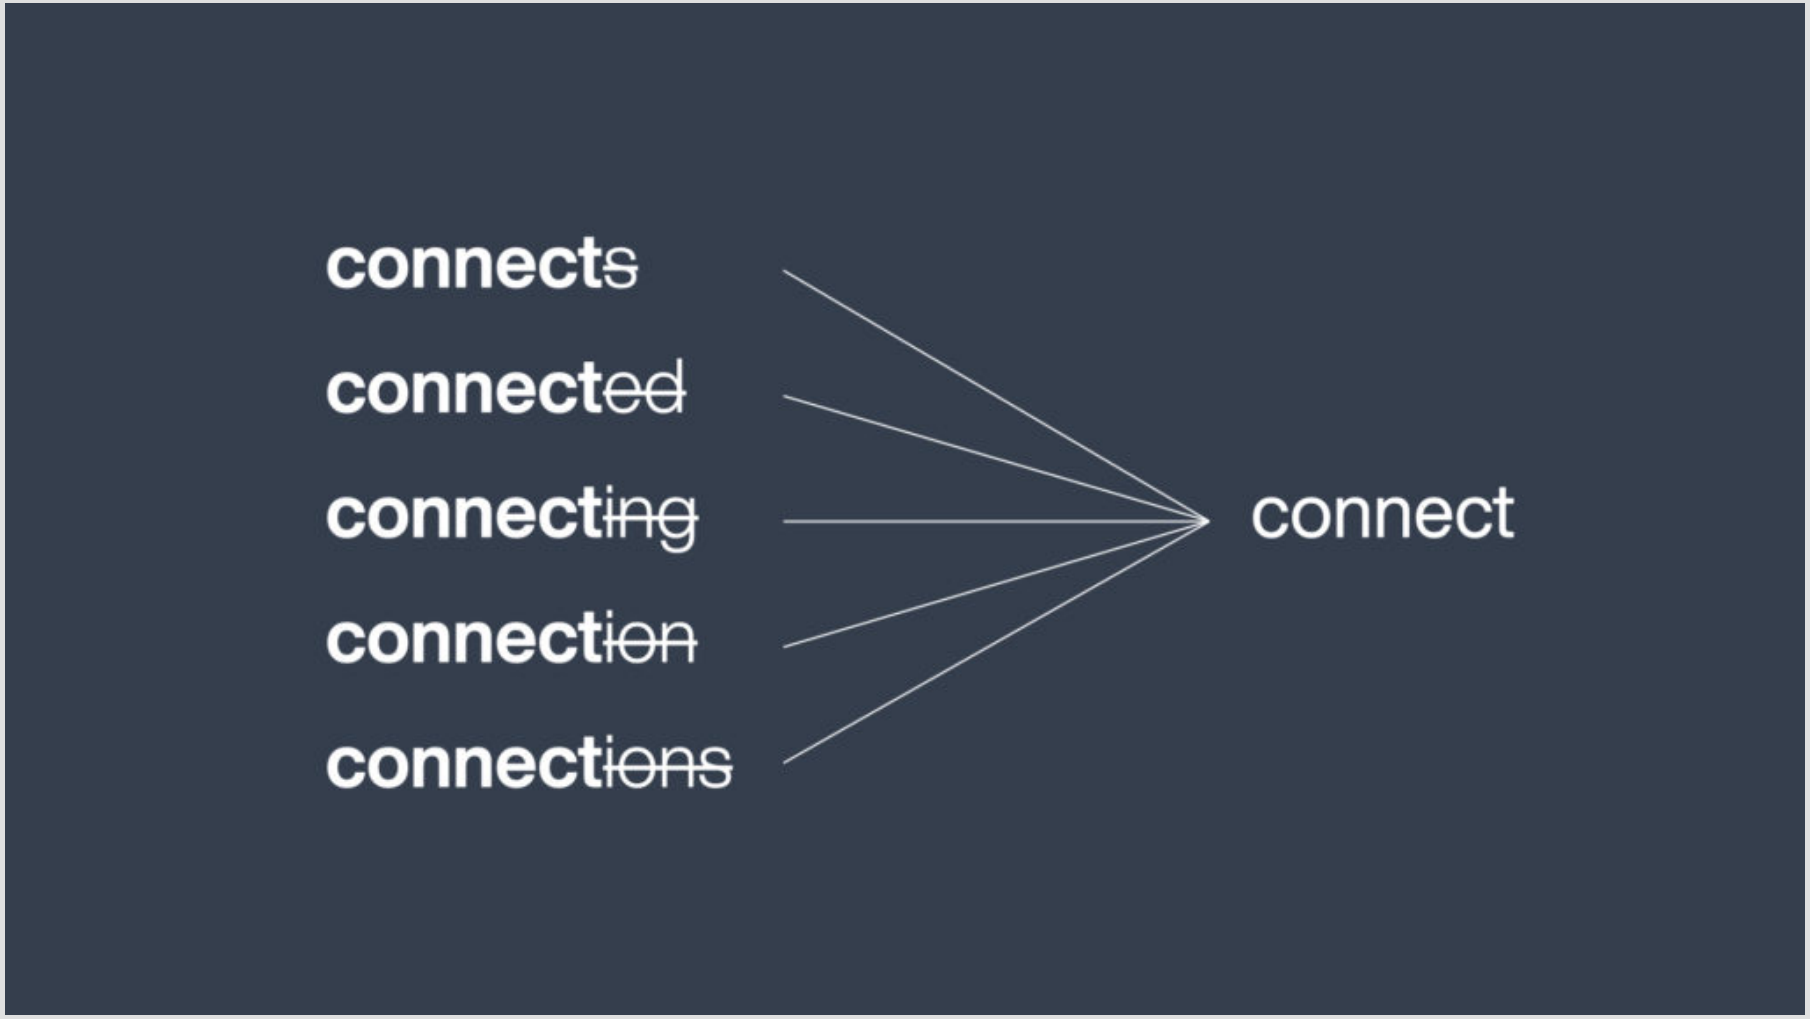

Early tasks: automatic summarization

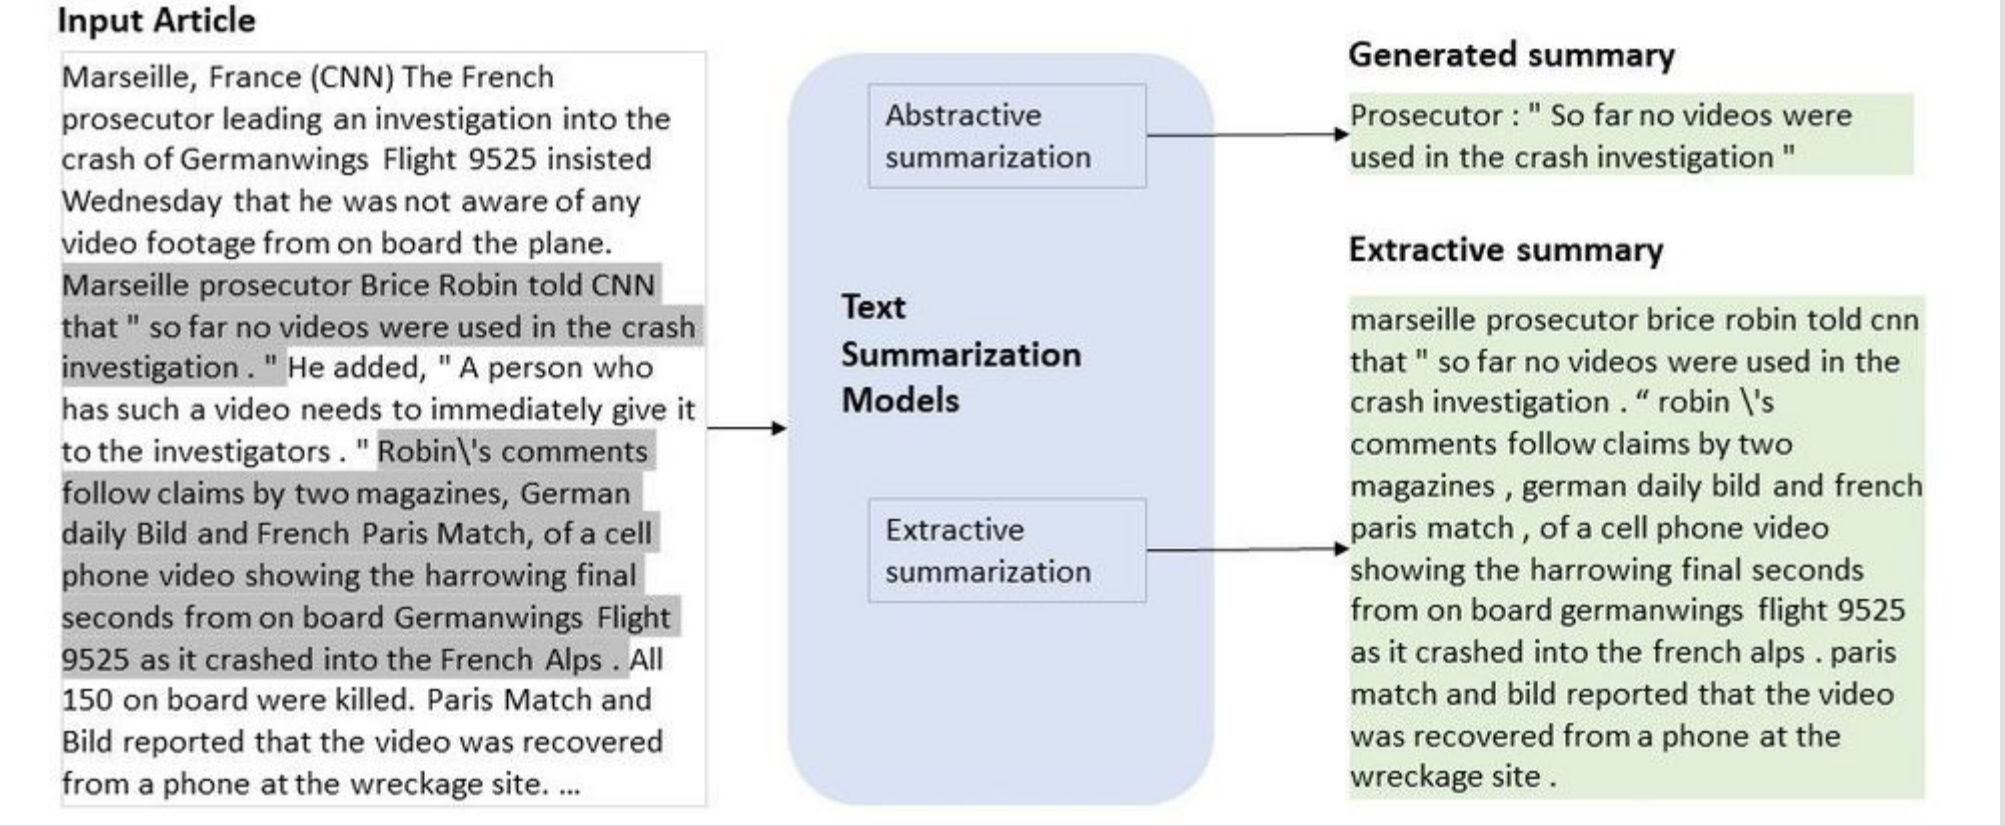

But how can we get at the *meaning* of natural language?

Text embedding models

- Preprocess some text to make a "training corpus"
- Train a model to parse the documents in the corpus
- Goal: generate "feature vectors" (for words, phrases, documents, etc.) that capture semantic properties of the text

Latent semantic analysis

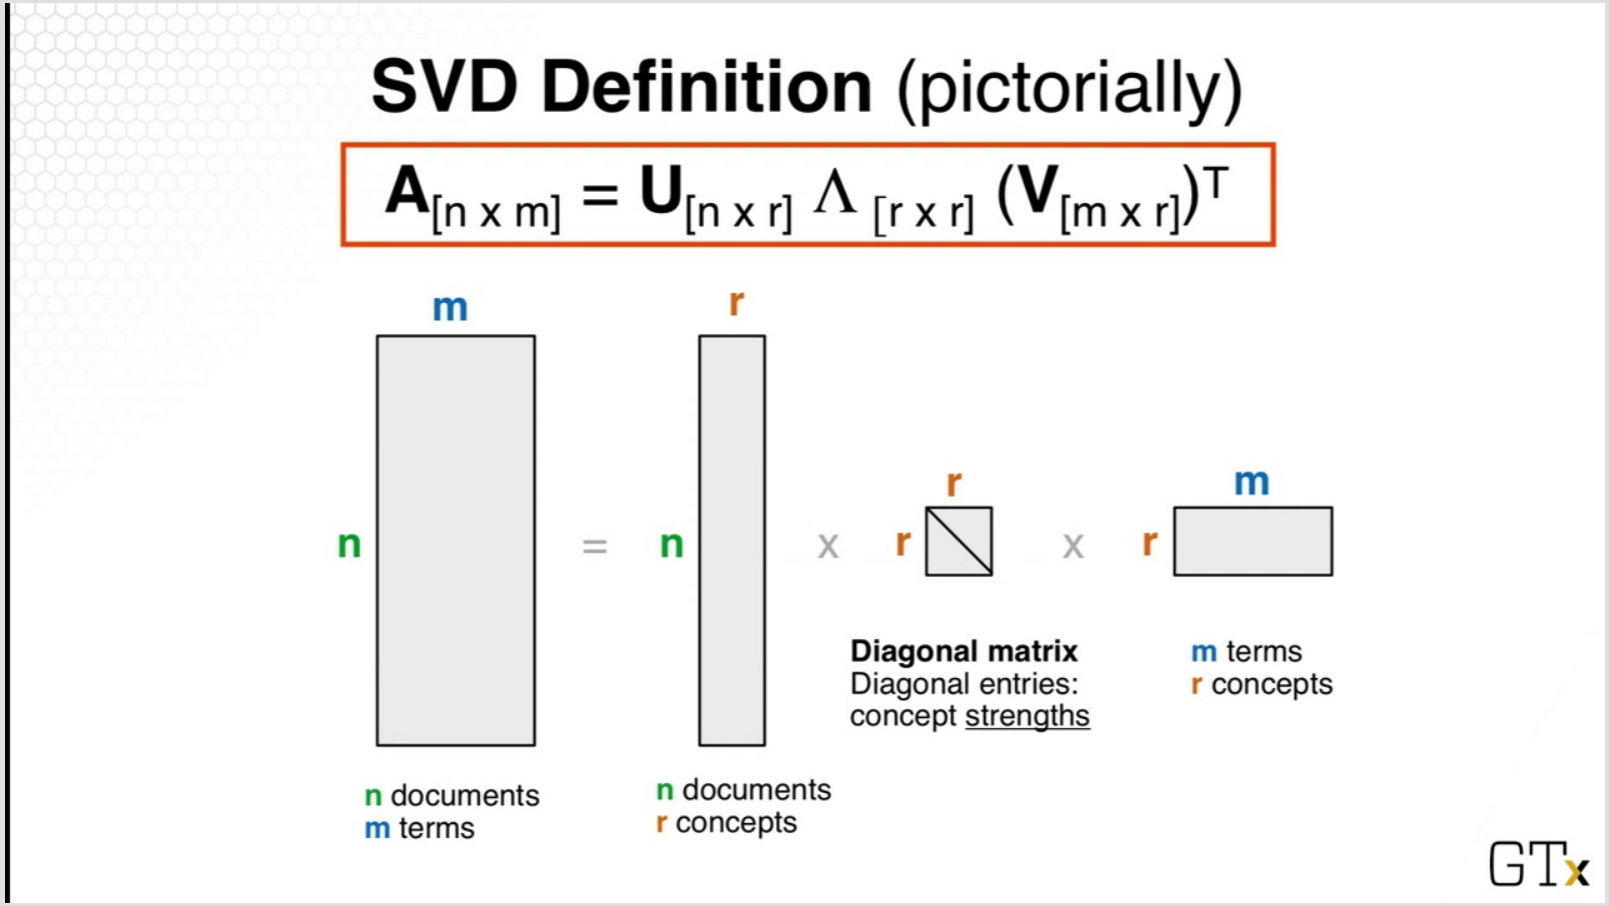

Latent Dirichlet Allocation (LDA)

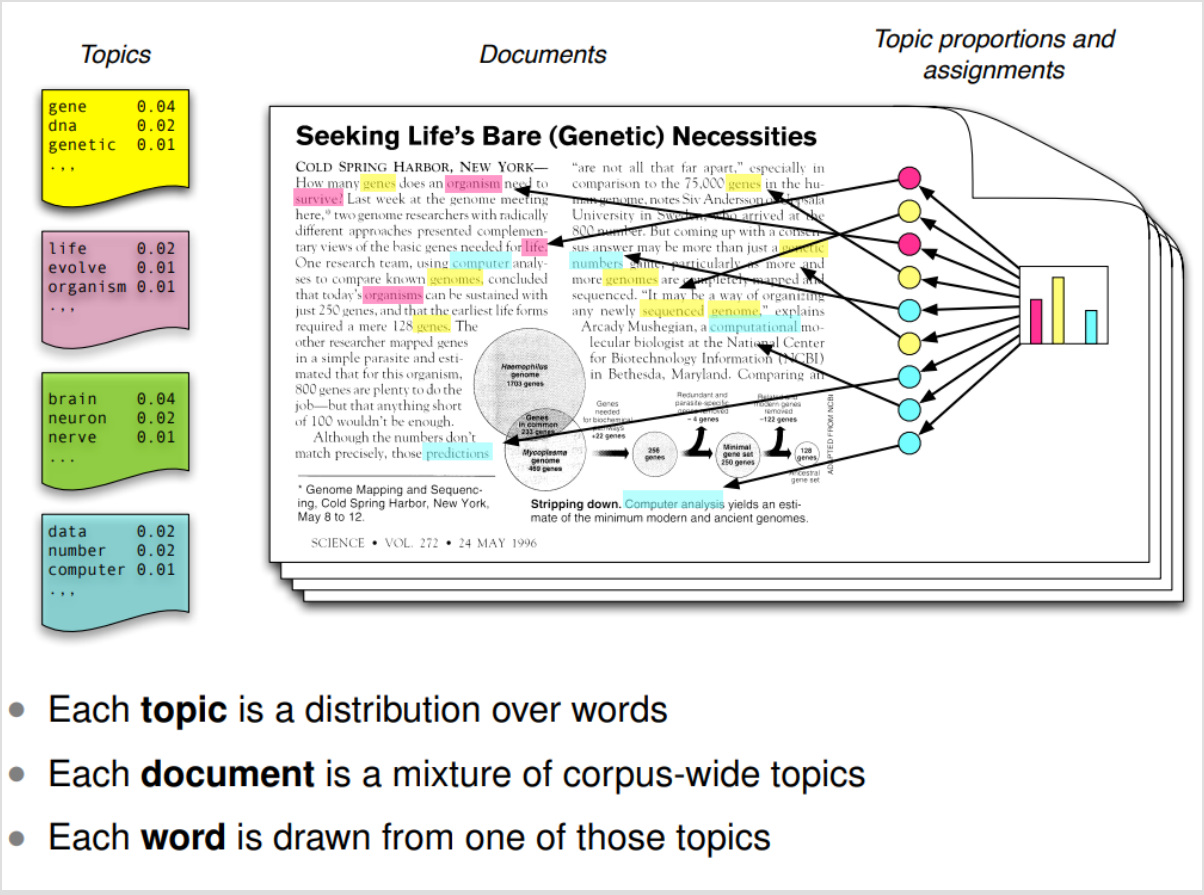

- Each **topic** is a distribution over words
- Each **document** is a mixture of corpus-wide topics
- Each **word** is drawn from one of those topics

Word2vec

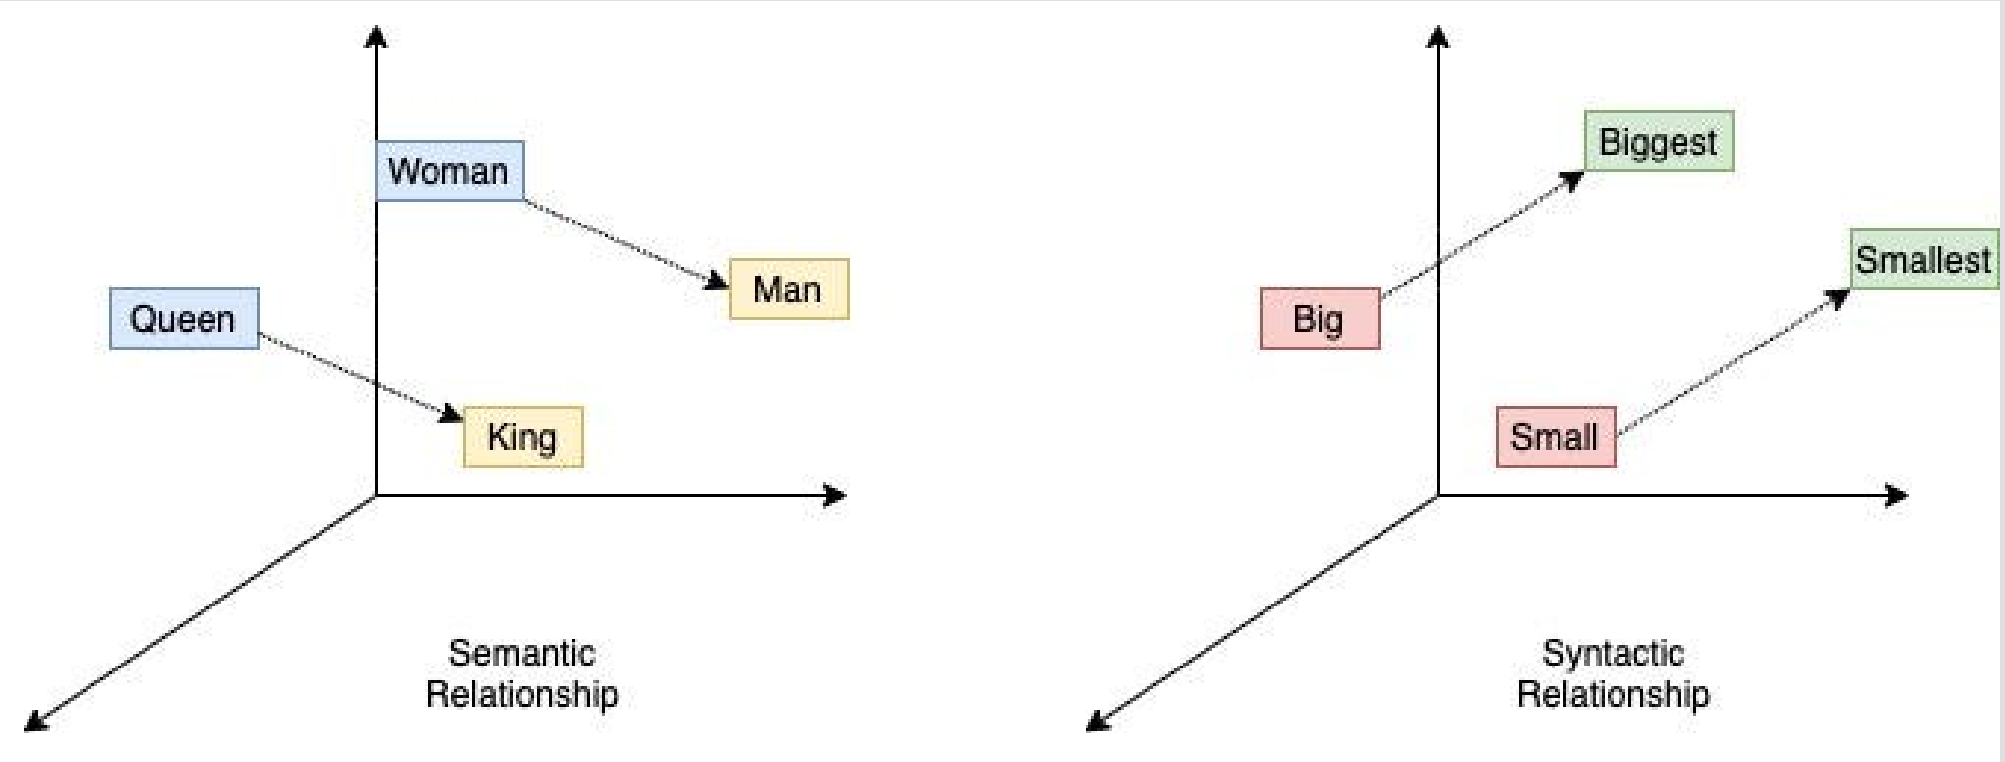
Woman
Man
Queen
King
Semantic
Relationship
Biggest
Smallest
Big
Small
Syntactic
Relationship

Consider the following phrases:

"My dog ate my homework"

vs.

"My homework ate my dog"

Bag of words vs. context-sensitive models

- BoW models (e.g., LSA, LDA, word2vec) don’t care about word order
- C-S models "care" about *context* and *grammar* by picking up on word order effects

Universal Sentence Encoder

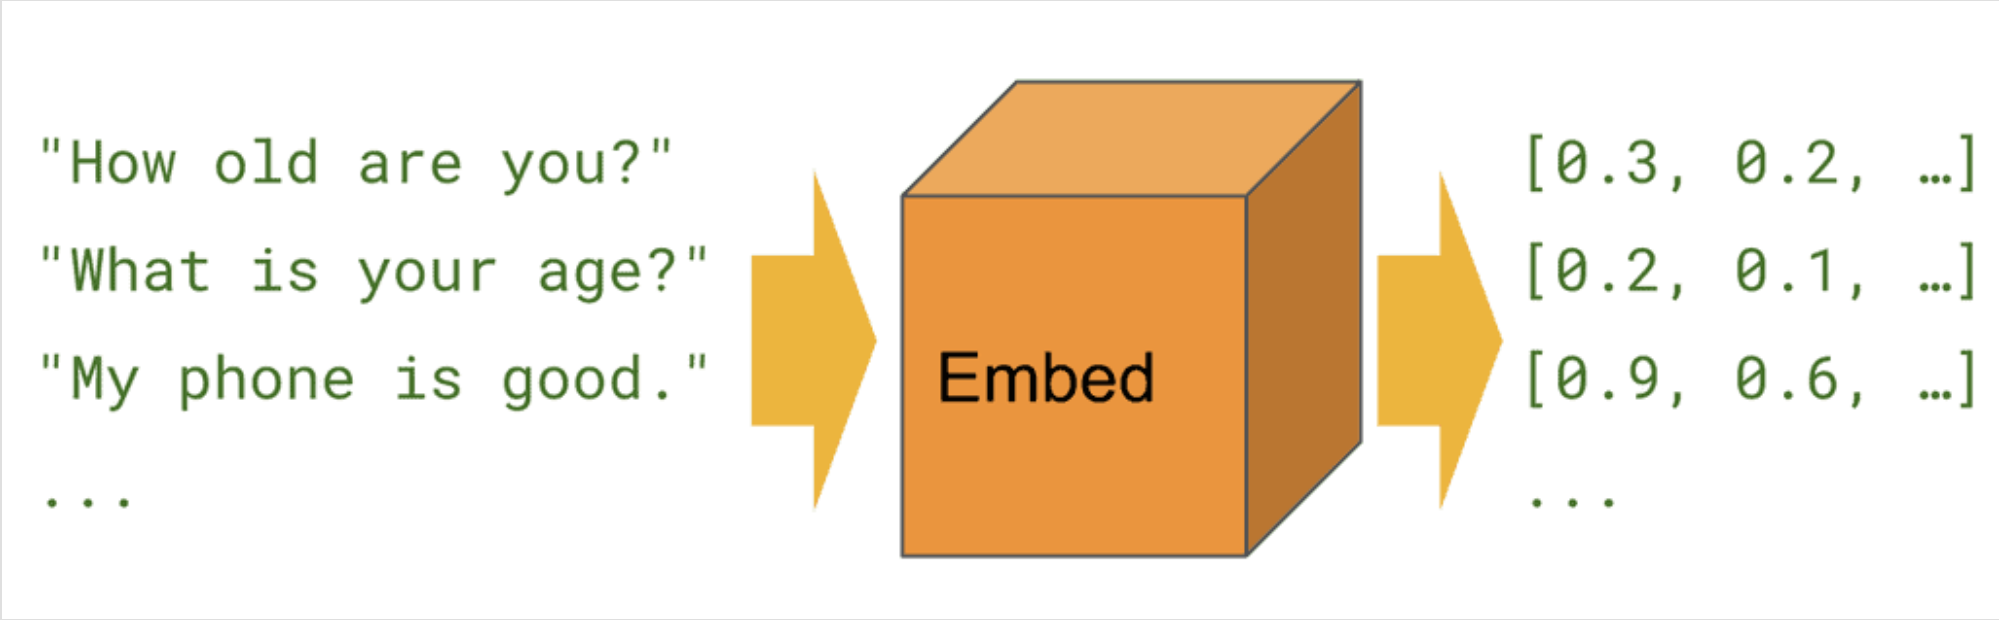

Transformers

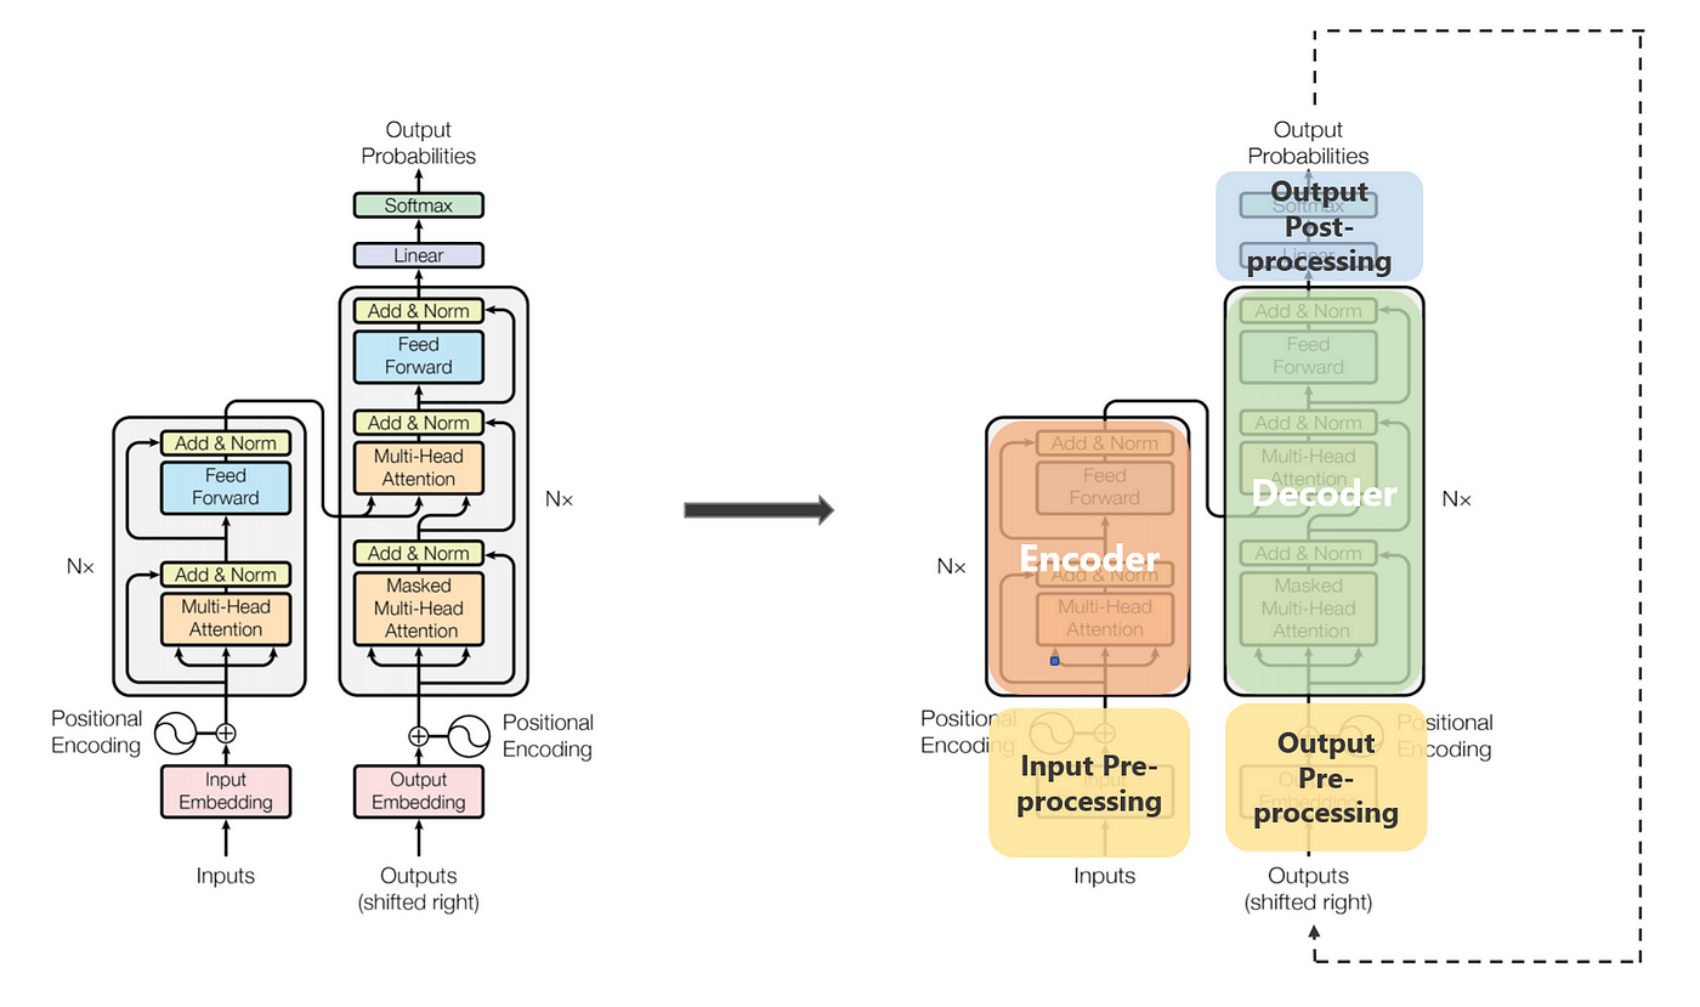

Transfomers

1. *Tokenize* text into smaller units (words/sub-words)
2. *Embed* tokens (ignoring order)
3. Update the embeddings to include *position* information (add features)
4. Update the embeddings using the *encoder and decoder* layers:
   a. The encoder processes the inputs
   b. The decoder generates outputs

Transformers

- Initially used for "sequence-to-sequence" tasks (e.g., translation)
- Now the basis of most state-of-the-art NLP models

Generative Pretrained Transformer (GPT)

- Variant of transformers
- "Regular" transformers have both an encoder and decoder; GPT only has the decoder part
- The goal is to predict the next token in a sequence, given the previous context
- The model is "pre-trained" on a large corpus of text to learn which tokens are likely to follow which other tokens
- The model can then be "fine-tuned" on specific tasks (e.g., Q&A, following instructions, coding tasks, etc.)
- Goal: text completion, translation, summarization, writing code, etc.

NLP in practice

- Natural Language Toolkit (NLTK) implements lots of basic text processing tasks like tokenization, lemmatization, part of speech tagging, etc.
- Scikit-learn has some basic models, like LDA
- For fancier models the best place to look is Hugging Face

NLTK

Implements lots of fundemantal "traditional" computational linguistics tasks

```
import nltk

sentence = """At eight o'clock on Thursday morning
... Arthur didn't feel very good."""
tokens = nltk.word_tokenize(sentence)

tagged = nltk.pos_tag(tokens)
```

Hugging Face

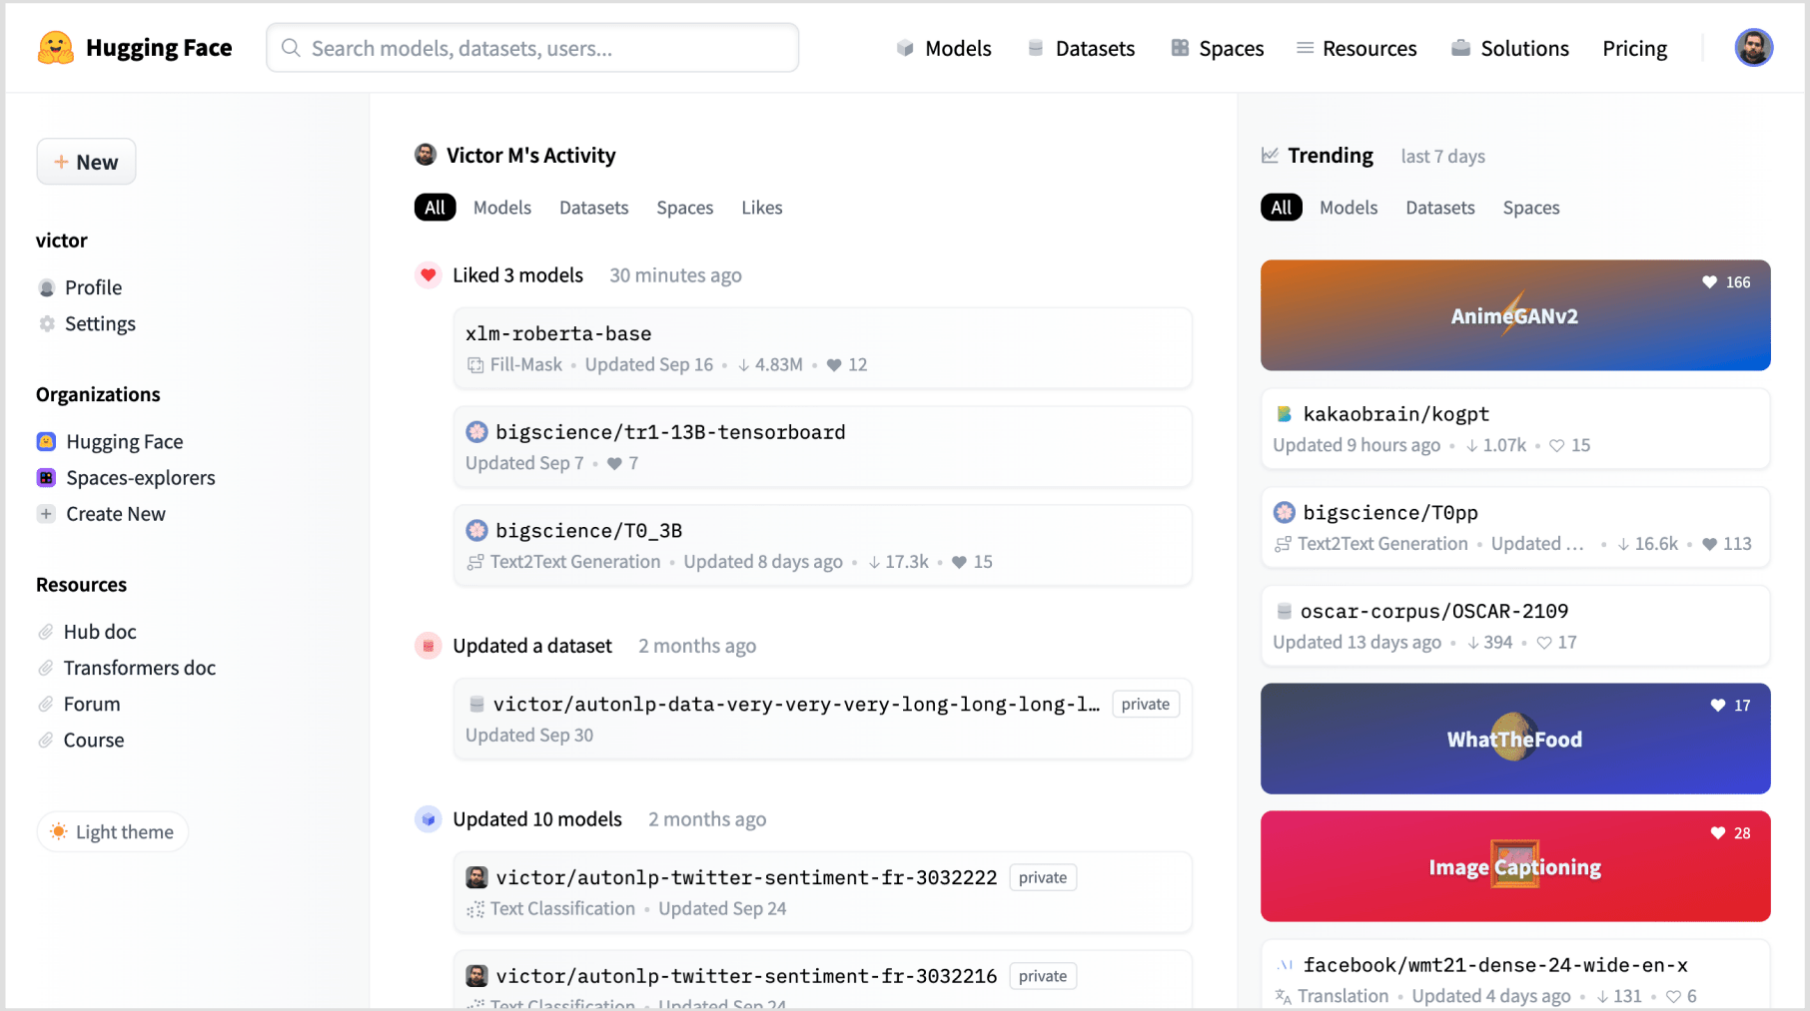

`transformers` library (by Hugging Face)

Direct interactions with text and models

```
from transformers import AutoTokenizer, AutoModelForMaskedLM

tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
model = AutoModelForMaskedLM.from_pretrained("bert-base-uncased")
```

`pydata-wrangler` library

- Wrapper for `scikit-learn` and Hugging Face models, focused on computing text embeddings

```
import datawrangler as dw

bert = {'model': 'TransformerDocumentEmbeddings',
        'args': ['bert-base-uncased'],
        'kwargs': {}}
bert_embeddings = dw.wrangle(my_text,
                             text_kwargs={'model': bert})
```

- Lots of other useful tools for working with data (supports numpy, pandas, text, images, and more)
- Essentially the core function (`dw.wrangle`) turns messy data into Pandas DataFrames

Interactive agents (ChatBots)

- Early implementations (e.g., ELIZA) used clever string manipulation hacks

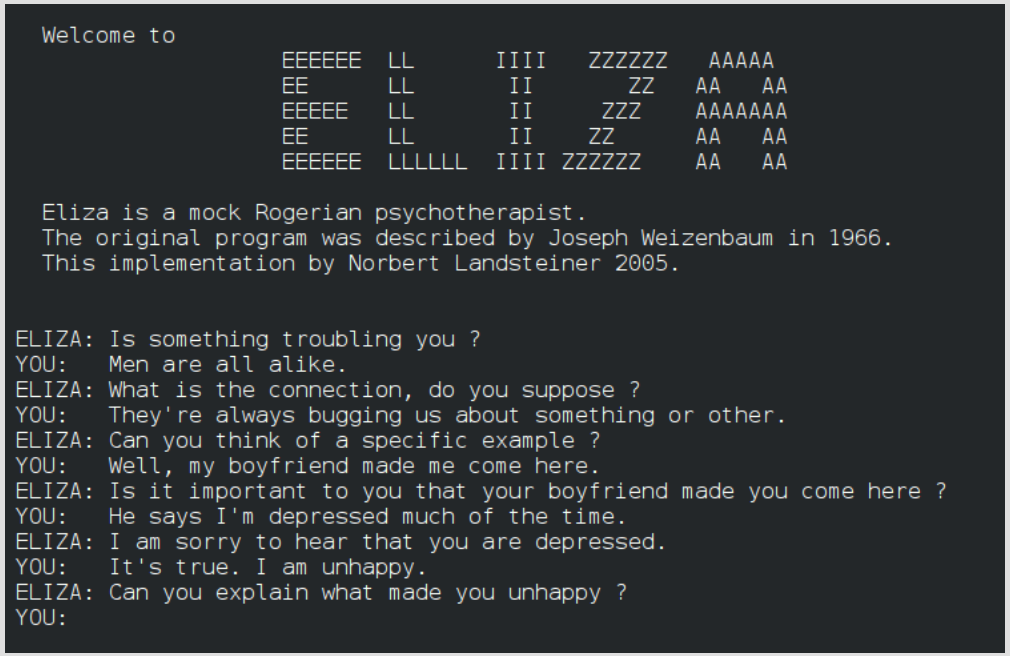

- Modern chatbots: predictive text models (text embeddings, positional coding, etc.)

`langchain` library

- Provides some great tools for interfacing with a wide variety of language models from Hugging Face, OpenAI, Meta, and more
- Infrastructure for chaining together prompts, tasks, models, etc.

Summary

- Text embeddings: bag of words vs. context-sensitive models
- Chatbots: string manipulations vs. predictive text models
- Suggested libraries: `scikit-learn`, `pydata-wrangler`, `transformers`, and `langchain`